

General Remarks

$d(x, y) = \ x - y\ $,	$\text{dist}(x, y) = 1 - \text{sim}(x, y)$
l_2 -euclidean distance	$\sqrt{\sum_{i=1}^d (x_i - y_i)^2}$
l_1 -manhattan distance	$\sum_{i=1}^d x_i - y_i $
l_p -distance	$(\sum_{i=1}^d x_i - y_i ^p)^{1/p}$
l_∞ -distance	$\max_i x_i - y_i $
Mahalanobis norm	$\ w\ _G^2 = \ Gw\ _2^2$
Cosine-Similarity	$\cos \frac{x^T y}{\ x\ _2 \ y\ _2}$
Jaccard-Distance	$1 - \text{sim}(A, B) = 1 - \frac{ A \cap B }{ A \cup B }$

Function Properties

Concave function

$$f(a + s) - f(a) \geq f(b + s) - f(b) \quad \forall a \leq b, s > 0$$

Convex functions A function $f : S \rightarrow \mathbb{R}$, $S \subseteq \mathbb{R}^d$, is called convex if $\forall x, x' \in S, \lambda \in [0, 1]$ it holds that $\lambda f(x) + (1 - \lambda)f(x') \geq f(\lambda x + (1 - \lambda)x')$

H-strongly convex

$$f(x') \geq f(x) + \nabla f(x)^T (x' - x) + \frac{H}{2} \|x' - x\|^2, H > 0$$

1-D: f is H-sc $\Leftarrow f''(x) \geq H, \forall x$.

d-D: f is H-sc $\Leftarrow \lambda_{\min}(\nabla^2 f(x)) \geq H, \forall x$.

Subgradients Given a convex not necessarily differentiable function f , a subgradient $g_x \in \nabla f(x)$ is the slope of a linear lower bound of f , tight at x , that is $\forall x' \in S : f(x') \geq f(x) + g_x^T (x' - x)$

Locality Sensitive Hashing

Near-duplicate detection

$$\{(x, y) \in X \times X : x \neq y, d(x, y) \leq \epsilon\}$$

(r, ϵ) -neighbour search Find all points with distance $\leq r$ and no points with distance $> (1 + \epsilon)r$ from query q . Pick $(r, (1 + \epsilon) \cdot r, p, q)$ -sensitive family and boost.

Min-hashing $h(C) = h_\pi(C) = \min_{i: C(i)=1} \pi(i)$

$$\pi(i) = h_{a,b}(i) = ((a \cdot i + b) \bmod p) \bmod N,$$

p prime (fixed) $> N$, N number of documents

- 1: **for** each column c **do**
- 2: **for** each row r **do**
- 3: **if** c has 1 in row r **then**
- 4: **for** each hash fn h_i **do**
- 5: $M_{i,c} \leftarrow \min\{h_i(r), M_{i,c}\}$

Band-hashing Signature matrix into b bands of r hash fns, per-column into b hash tables. If any h-table has a collision, report candidate pair. s^r prob of col on band j : $P(\text{col in } j \geq 1 \text{ band}) = 1 - (1 - s^r)^b$

(d_1, d_2, p_1, p_2) -sensitivity Assume $d_1 < d_2, p_1 > p_2$.
 $\forall x, y \in S : d(x, y) \leq d_1 \Rightarrow \Pr[h(x) = h(y)] \geq p_1$
 $\forall x, y \in S : d(x, y) \geq d_2 \Rightarrow \Pr[h(x) = h(y)] \leq p_2$

r-way AND $h(x) = h(y) \iff \forall i : h_i(x) = h_i(y)$

(d_1, d_2, p_1^r, p_2^r) - big r , more FN

b-way OR $h(x) = h(y) \iff \exists i : h_i(x) = h_i(y)$

$(d_1, d_2, 1 - (1 - p_1)^b, 1 - (1 - p_2)^b)$ - big b , more FP

AND-OR cascade $(d_1, d_2, 1 - (1 - p_1^r)^b, 1 - (1 - p_2^r)^b)$

OR-AND cascade

$(d_1, d_2, (1 - (1 - p_1)^b)^r, (1 - (1 - p_2)^b)^r)$

Hash Functions

Euclidean distance $h_{w,b}(x) = \lfloor \frac{w^T x - b}{a} \rfloor$ where

$$w \leftarrow \frac{w}{\|w\|_2}, w \sim \mathcal{N}(0, I), w_i \sim \mathcal{N}(0, 1),$$

$b \sim \text{Unif}([0, a])$, yields $(a/2, 2a, 1/2, 1/3)$ -sensitive

Cosine distance $\mathcal{H} = \{h(v) = \text{sgn}(w^T v)\}$ where

$$w \sim \text{Unif}\{x \in \mathbb{R}^n : \|x\|_2 = 1\}$$

$$\Pr(h_u(x) = h_v(y)) = 1 - \Theta_{x,y}/\pi$$

Support Vector Machines

SVM SVM = Max margin linear classifier

$$\min_{w, \xi \geq 0} \frac{1}{2} w^T w + C \sum_{i=1}^n \xi_i \text{ s. t. } y_i w^T x_i \geq 1 - \xi_i \quad \forall i$$

Support vectors (SV) are all data points on the margin and data points with non-zero slack

Regularized hinge loss formulation $C = 1/\lambda$

$$\min_w \lambda w^T w + C \sum_i \max(0, 1 - y_i w^T x_i)$$

Norm-constrained hinge loss minimization

$$\min_w \sum_i \max(0, 1 - y_i w^T x_i) \text{ s.t. } \|w\|_2 \leq \frac{1}{\sqrt{\lambda}}$$

Strongly convex formulation

$$\min_w \frac{1}{T} \sum_{t=1}^T \left(\frac{\lambda}{2} \|w\|_2^2 + \max(0, 1 - y_t w^T x_t) \right)$$

$$\text{s.t. } \|w\|_2 \leq \frac{1}{\sqrt{\lambda}}$$

$$\text{Proj}_S^{(2)}(w) = w \cdot \min \left(1, \frac{1/\sqrt{\lambda}}{\|w\|} \right)$$

$$\text{Proj}_S^{(\infty)}(w) = \max((1, \dots, 1)^T, w)$$

Small C , Big λ : Greater margin, more misclassification

Kernels

Dual SVM Formulation

$$\max_\alpha \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j x_i^T x_j$$

$$\text{s.t. } 0 \leq \alpha_i \leq C$$

$$\Rightarrow \text{optimal } w : w^* = \sum_i \alpha_i^* y_i x_i = \sum_{i \in \text{SV}} \alpha_i^* y_i x_i$$

Kernel trick Substitute inner product $x_i^T x_j$ in dual formulation and in classification function with $k(x_i, x_j) = \phi(x_i)^T \phi(x_j)$, where $\phi(\cdot) : \mathbb{R}^d \rightarrow \mathbb{R}^{>d}$

Kernel functions A kernel is function $k : X \times X \rightarrow \mathbb{R}$:

1. Symmetry: $\forall x, x' \in X : k(x, x') = k(x', x)$
2. PSD: $\forall n \in \mathbb{N}$, any set $S = \{x_1, \dots, x_n\} \subseteq X$, the Gram matrix is PSD.

Random Features (Inverse Kernel Trick)

Shift-invariant kernel $k(x, y) = k'(x - y)$. Then the kernel has Fourier transform, such that:

$$k(x - y) = \int_{\mathbb{R}^d} p(w) \cdot e^{i w^T (x - y)} dw$$

where $p(w)$ is the Fourier transformation, i.e. we map $k(s)$ to another function $p(w)$.

Random fourier features (prerequisites) Interpret kernel as expectation $k(x - y) =$

$$\int_{\mathbb{R}^d} p(w) \cdot \underbrace{e^{i w^T (x - y)}}_{g(w)} dw = \mathbb{E}_{w,b} [z_{w,b}(x) z_{w,b}(y)]$$

$$\text{where } z_{w,b}(x) = \sqrt{2} \cos(w^T x + b),$$

$$b \sim U([0, 2\pi]), w \sim p(w)$$

Random fourier features (kernel approximation)

1. $w_i \sim p, b_i \sim U([0, 2\pi])$ for $i = 1, \dots, m$; iid
2. $z(x) = [z_{w_1, b_1}(x), \dots, z_{w_m, b_m}(x)] / \sqrt{m}$
3. $z(x)^T z(y) = \frac{1}{m} \sum_{i=1}^m z_{w_i, b_i}(x) \cdot z_{w_i, b_i}(y)$
4. If $m \rightarrow \infty$, then (almost surely)
 $z(x)^T z(y) \rightarrow \mathbb{E}_{w,b} [z_{w,b}(x) \cdot z_{w,b}(y)] = k(x - y)$

Online Convex Programming

Regret $R_T = (\sum_{t=1}^T f_t(w_t)) - \min_{w \in S} \sum_{t=1}^T f_t(w)$

No-regret $\lim_{T \rightarrow \infty} \frac{R_T}{T} \rightarrow 0$

Online convex programming (OCP) If $y_t w_t^T x_t < 1 : w_{t+1} = \text{Proj}_S(w_t - \eta_t \nabla f_t(w_t))$

$$\text{Proj}_S(w) = \arg \min_{w' \in S} \|w' - w\|_2 = \min \left(w, \frac{w}{\lambda \|w\|_2} \right)$$

Regret for OCP

$$\frac{R_T}{T} \leq \frac{1}{\sqrt{T}} [\|w_0 - w^*\|_2^2 + \|\nabla f\|_2^2]$$

where $\|\nabla f\|_2^2 = \sup_{w \in S, t \in \{1, \dots, T\}} \|\nabla f(w)\|_2^2$

Parallel stochastic gradient descent 1. Split data into k subsets, k = number of machines, want $k = O(1/\lambda)$

2. Each machine produces w_i on its subset

3. After T iterations, compute $w = \frac{1}{k} \sum_{i=1}^k w_i$

PEGASOS Online SVM: H-sc. lossfn. Minibatch + reg

Active Learning

Uncertainty sampling Repeat until all labels inferred:

1. Assign uncertainty score $U_t(x)$ to each unlabeled data point:
 $U_t(x) = U(x|_{x_{1:t-1}, y_{1:t-1}})$
2. Greedily pick the most uncertain point and request label $x_t = \arg \max_x U_t(x)$ and retrain classifier

For SVM: $U_t(x) = \frac{1}{|w_t^T x|}$

Cost to pick m labels: $m \cdot n \cdot d + m \cdot C(m)$

n = number of data points, d = dimensions,

$C(m)$ = cost to train classifier

Hashing a hyperplane query Draw $u, v \sim \mathcal{N}(0, I)$.

Then resulting two-bit hash is:

$$h_{u,v}(a, b) = \left[\text{sign}(u^T a), \text{sign}(v^T b) \right]$$

Define the hash family:

$$h_{\mathcal{H}}(z) = \begin{cases} h_{u,v}(z, z) & \text{if } z \text{ is a database point vector} \\ h_{u,v}(z, -z) & \text{if } z \text{ is a query hyperplane vector} \end{cases}$$

Version space Set of all classifiers consistent with the data: $V(D) = \{w : \forall (x, y) \in D : \text{sign}(w^T x) = y\}$

Relevant version space $\hat{V}(D; U)$ describes all possible labelings h of all unlabeled data U that are still possible under some model w , or,

$$\hat{V}(D; U) = \{h : U \rightarrow \{+1, -1\} : \exists w \in V(D) \forall x \in U : \text{sign}(w^T x) = h(x)\}$$

Generalized Binary Search (GBS)

- 1: Start with $D = \emptyset$
- 2: **while** $|\hat{V}(D; U)| > 1$ **do**
- 3: **for each** unlabeled example x in U **do**
- 4: $v^+(x) = |\hat{V}(D \cup \{(x, +)\}; U)|$
- 5: $v^-(x) = |\hat{V}(D \cup \{(x, -)\}; U)|$
- 6: ▷ number of labelings left if x is $-/+$
- 7: Pick $x^* = \arg \min_x \max(v^-(x), v^+(x))$
- 8: $D = D \cup \{x^*\}$

Decision rules for GBS SVM, $m \sim \text{margin} \sim \frac{1}{\|w\|}$

Max-min margin $\max_x \min(m^+(x), m^-(x))$

Ratio margin $\max_x \min\left(\frac{m^+(x)}{m^-(x)}, \frac{m^-(x)}{m^+(x)}\right)$

Clustering

K-Means

Cost Function $L(\mu) = L(\mu_1, \dots, \mu_k) =$

$$\sum_{i=1}^N \underbrace{\min_{j \in \{1, \dots, k\}} \|x_i - \mu_j\|_2^2}_{d(\mu, x_i)}$$

Objective $\mu^* = \arg \min_{\mu} L(\mu)$

Algorithm Until convergence:

- 1: Assign each point x_i to closest center

$$z_i \leftarrow \arg \min_{j \in \{1, \dots, l\}} \|x_i - \mu_j^{(t-1)}\|_2^2$$

- 2: Update center as mean of assigned data points

$$\mu_j^{(t)} \leftarrow \frac{1}{n_j} \sum_{i: z_i=j} x_i$$

Online k-means algorithm

$$\frac{dd(\mu, x_t)}{d\mu_j} = \begin{cases} 0 & \text{if } j \notin \arg \min_i \|\mu_i - x_t\|^2 \\ 2(u_j - x_t) & \text{else} \end{cases}$$

1. Initialize centers randomly
2. For $t = 1 : N$
 - Find $c = \arg \min \|\mu_j - x_t\|_2$
 - $\mu_c = \mu_c + \eta_t(x_t - \mu_c)$
3. For convergence: $\sum_t \eta_t = \infty$ and $\sum_t \eta_t^2 < \infty$, e.g. $\eta_t = \frac{c}{t}$.

Coresets

Key idea Replace many points by one weighted representative, thus, obtain C

$$L_k(u; C) = \sum_{(w, x) \in C} w \cdot \min_j \|u_j - x\|_2^2$$

(k, ϵ) -coreset C is called a (k, ϵ) -coreset for D , if for all μ : $(1 - \epsilon)L_k(\mu; D) \leq L_k(\mu; C) \leq (1 + \epsilon)L_k(\mu; D)$

Operations Merge: union of two (k, ϵ) -cses is a (k, ϵ) -cset

Compress: (k, δ) -cset of a (k, ϵ) -cset of D is a $(k, \epsilon + \delta + \epsilon\delta)$ -cset of D

Construction D^2 -sampling: iteratively build $B = \emptyset$ by sampling $\propto \frac{d(x, B)^2}{\sum_{x' \in X} d(x', B)^2}$

Importance sampling:

$$\propto \frac{\alpha d(x, B)^2}{c_{\Phi}} + \frac{2\alpha \sum_{x' \in B_x} d(x', B)^2}{|B_x| c_{\Phi} h_i} + \frac{4|X|}{|B_x|} \text{ where}$$

$c_{\Phi} = \frac{1}{|X|} \sum_{x \in X} d(x, B)^2$
 B_x = pts in X that belong to same cluster as x

Bandits

Regret $R_T = \sum_{t=1}^T (\mu^* - \mu_{i_t})$, i_t chosen arm at time t
 ϵ -greedy The algorithm goes as follows:

1. Set $\epsilon_t = \mathcal{O}(\frac{1}{t})$
2. With probability ϵ_t : explore by picking uniformly at random
3. With probability $1 - \epsilon_t$: exploit by picking arm with highest empirical mean

Regret: $R_T = \mathcal{O}(k \log(T))$

UCB1 & LinUCB

Hoeffding's inequality

$$\Pr(|\mu - \frac{1}{m} \sum_{t=1}^m X_t| \geq b) \leq 2 \cdot \exp(-2b^2 m)$$

Confidence bound Want: Hoeffding $\leq \delta$

$$\Rightarrow b = \sqrt{\frac{1}{2m} \ln \frac{2}{\delta}}$$

UCB/Mean update $UCB(i) = \hat{\mu}_i + \sqrt{\frac{2 \ln t}{\eta_i}}$

$$j = \arg \max_i UCB(i)$$

$$\hat{\mu}_j = \hat{\mu}_j + \frac{1}{\eta_j} (y_t - \hat{u}_j)$$

Contextual bandits Reward is now $y_t = f(x_t, z_t) + \epsilon_t$ with z_t user features. For us: $f(x_i, z_t) = w_{x_i}^T z_t$

LinUCB (disjoint) Ridge regression on z_t :

$$\hat{w}_i = (D_i^T D_i + I)^{-1} D_i^T y_i.$$

$$|\hat{w}_i^T z_t - w_i^T z_t| \leq \alpha \sqrt{z_t^T (D_i^T D_i)^{-1} z_t} \text{ with probability } 1 - \delta \text{ if } \alpha = 1 + \sqrt{\log(2/\delta)/2}$$

$$R_T/T = \mathcal{O}(d \cdot d' \cdot \text{poly log } T / \sqrt{T})$$

Hybrid model Reward is now

$$y_t = w_{x_t}^T z_t + \beta^T \phi(x_t, z_t) + \epsilon_t$$

Rejection sampling First obtain data log through pure exploration, and then reiterate:

1. Get event $(x_t^{(1)}, \dots, x_t^{(k)}, z_t, a_t, y_t)$ from log
2. Use algorithm that is testing to pick a'_t :
 - If $a'_t = a_t \Rightarrow$ Feed back reward y_t
 - Else ignore log line
3. Stop when T rewards have been fed back

Submodular Functions

Submodularity A function $F : 2^V \mapsto \mathbb{R}$ is called submodular iff for all $A \subseteq B, s \notin B$:

$$F(A \cup \{s\}) - F(A) \geq F(B \cup \{s\}) - F(B)$$

Union-intersection def. Same preconditions as above:

$$F(A) - F(A \cap B) \geq F(A \cup B) - F(B)$$

Closure properties F, F_i submodular on V ; $S, W \subseteq V$

Linear Combinations

$$F'(S) = \sum_i \lambda_i F_i(S), \lambda_i \geq 0$$

Restriction $F'(S) = F(S \cap W)$

Conditioning $F'(S) = F(S \cup W)$

Reflection $F'(S) = F(V \setminus S)$

Monotonic truncation F also monotonic,

$$F'(S) = \min(c, F(S)), c \in \mathbb{R}$$

Min/Max For $F_{1,2}(A)$, $\max\{F_1(A), F_2(A)\}$ or $\min\{F_1(A), F_2(A)\}$ **not** submodular in general.

Concavity $F(A) = g(|A|)$ where $g : \mathbb{N} \mapsto \mathbb{R}$, then F submodular iff g concave.

Lazy Greedy Optimizing submodular **monotonic** functions.

- 1: $A_0 = \emptyset$. Keep ordered list of marginal benefits Δ_i from previous iterations
- 2: **for** $i = 1 \dots k$: **do**
- 3: $\Delta_i = F(A_{i-1} \cup \{\text{top element}\}) - F(A_{i-1})$
- 4: **if** i is not top element anymore **then**
- 5: re-sort the list
- 6: $A_i = A_{i-1} \cup \{\text{top element}\}$